




Graph Output




Histogram of missing data



Pattern







Gender




Married




Dependents




Education




Self\_Employed




Credit\_History




Property\_Area



Loan\_Status


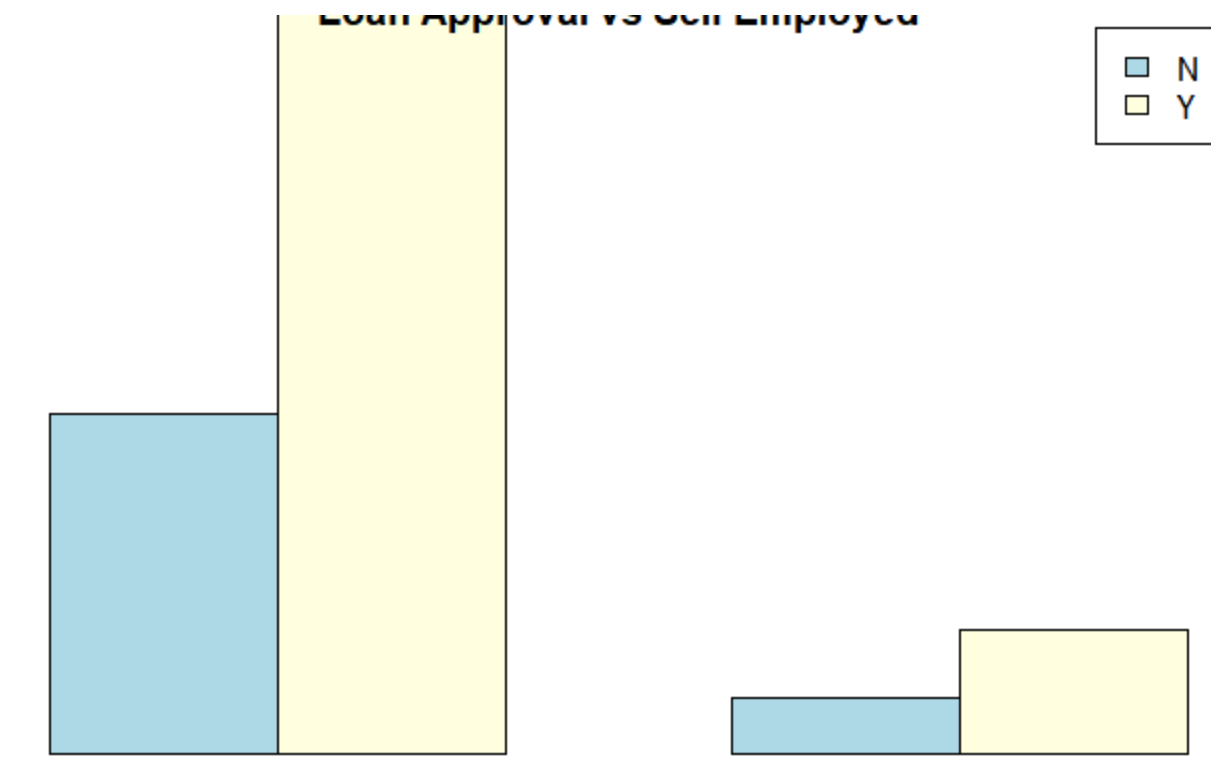


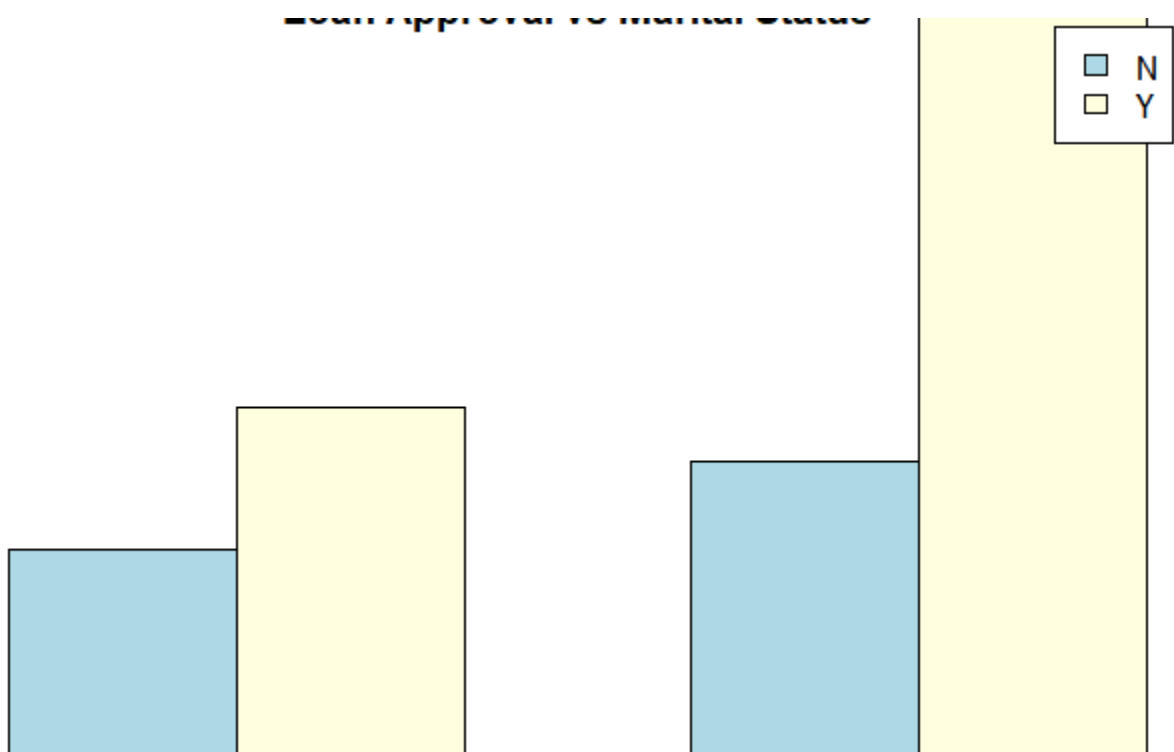
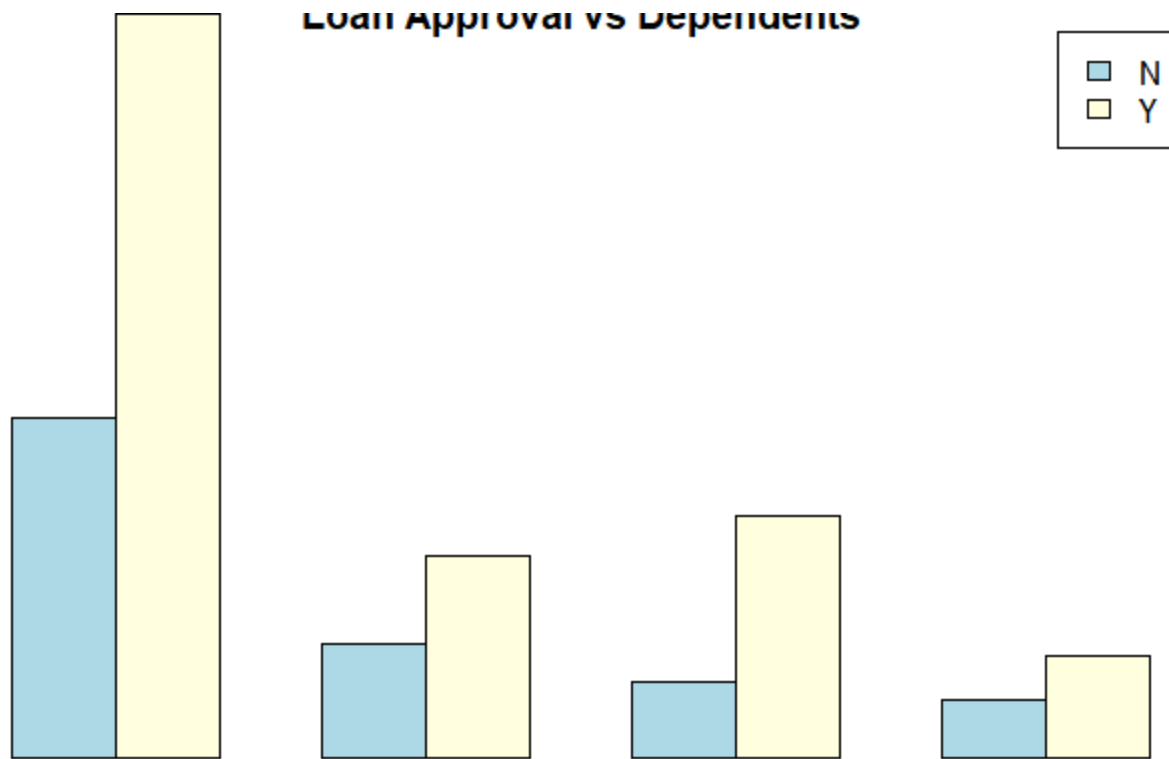
Loan Approval vs Credit history





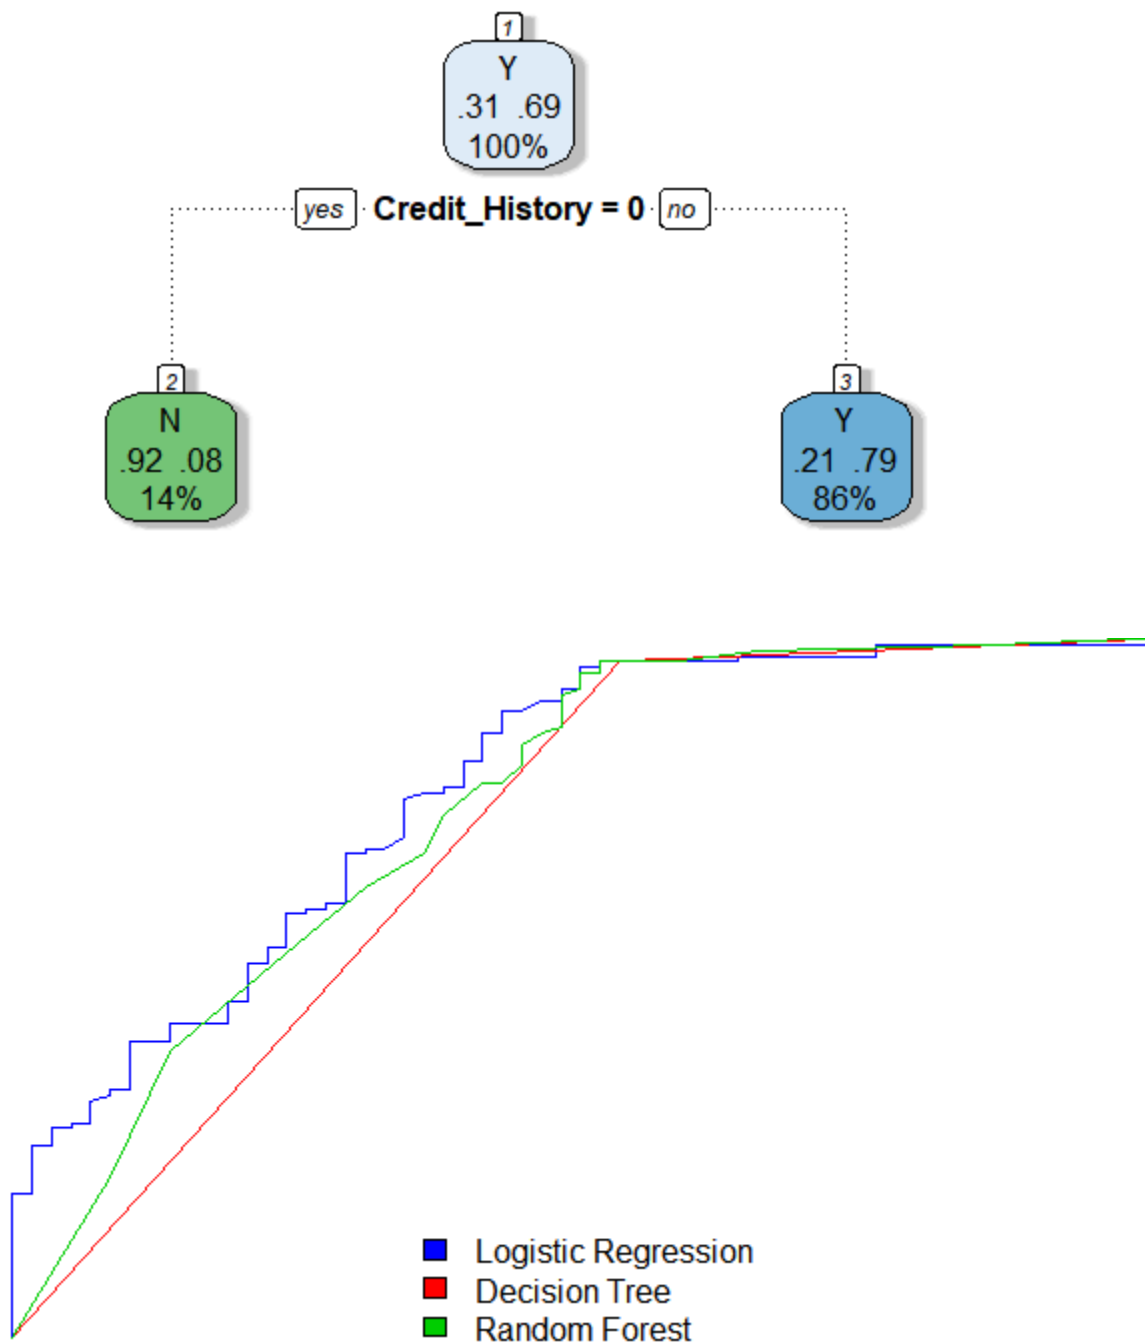
Loan Approval vs Loan Amount term











Code output

```
> set.seed(108)
> work_path = "D:/vaibhav/trend nxt/topgear/R Community/Predictive Model Based Logistic Regression-LoanData"
> setwd(work_path)
```

```

> # Some NAs are coded as empty strings , hence converting them to NA
> data <- read.csv("Loan_data.csv", header=T, na.strings=c("", " ", "NA"))
> original_data=data
> head(data)
  Loan_ID Gender Married Dependents Education Self_Employed ApplicantIncome
1 LP001002 Male No 0 Graduate No 58
2 LP001003 Male Yes 1 Graduate No 45
3 LP001005 Male Yes 0 Graduate Yes 30
4 LP001006 Male Yes 0 Not Graduate No 25
5 LP001008 Male No 0 Graduate No 60
6 LP001011 Male Yes 2 Graduate Yes 54

  CoapplicantIncome LoanAmount Loan_Amount_Term Credit_History Property_Area
1 0 NA 360 1 Urban
2 1508 128 360 1 Rural
3 0 66 360 1 Urban
4 2358 120 360 1 Urban
5 0 141 360 1 Urban
6 4196 267 360 1 Urban

  Loan_Status
1 Y
2 N
3 Y
4 Y
5 Y
6 Y
> summary(data)
  Loan_ID Gender Married Dependents Education Self_Employed
LP001002: 1 Female:112 No :213 0 :345 Graduate :480 No :50
LP001003: 1 Male :489 Yes :398 1 :102 Not Graduate:134 Yes : 8
LP001005: 1 NA's : 13 NA's: 3 2 :101 NA's: 3
LP001006: 1 3+ : 51
LP001008: 1 NA's: 15
LP001011: 1
(Other) :608
ApplicantIncome CoapplicantIncome LoanAmount Loan_Amount_Term
Min. : 150 Min. : 0 Min. : 9.0 Min. : 12
1st Qu.: 2878 1st Qu.: 0 1st Qu.:100.0 1st Qu.:360
Median : 3812 Median : 1188 Median :128.0 Median :360
Mean : 5403 Mean : 1621 Mean :146.4 Mean :342
3rd Qu.: 5795 3rd Qu.: 2297 3rd Qu.:168.0 3rd Qu.:360
Max. :81000 Max. :41667 Max. :700.0 Max. :480
NA's :22 NA's :14

Credit_History Property_Area Loan_Status
Min. :0.0000 Rural :179 N:192
1st Qu.:1.0000 Semiurban:233 Y:422
Median :1.0000 Urban :202
Mean :0.8422
3rd Qu.:1.0000
Max. :1.0000
NA's :50
> str(data)
'data.frame': 614 obs. of 13 variables:

```



```

$ Loan_ID      : Factor w/ 614 levels "LP001002","LP001003",...: 1 2 3 4
5 6 7 8 9 10 ...
$ Gender      : Factor w/ 2 levels "Female","Male": 2 2 2 2 2 2 2 2 2 2
...
$ Married     : Factor w/ 2 levels "No","Yes": 1 2 2 2 1 2 2 2 2 2 ...
$ Dependents  : Factor w/ 4 levels "0","1","2","3+": 1 2 1 1 1 3 1 4 3
2 ...
$ Education   : Factor w/ 2 levels "Graduate","Not Graduate": 1 1 1 2 1
1 2 1 1 1 ...
$ Self_Employed : Factor w/ 2 levels "No","Yes": 1 1 2 1 1 2 1 1 1 1 ...
$ ApplicantIncome : int  5849 4583 3000 2583 6000 5417 2333 3036 4006 12841
...
$ CoapplicantIncome: num  0 1508 0 2358 0 ...
$ LoanAmount      : int  NA 128 66 120 141 267 95 158 168 349 ...
$ Loan_Amount_Term : int  360 360 360 360 360 360 360 360 360 360 ...
$ Credit_History   : int  1 1 1 1 1 1 1 0 1 1 ...
$ Property_Area    : Factor w/ 3 levels "Rural","Semiurban",...: 3 1 3 3 3 3
3 2 3 2 ...
$ Loan_Status     : Factor w/ 2 levels "N","Y": 2 1 2 2 2 2 2 1 2 1 ...
> nrow(data)
[1] 614
> ncol(data)
[1] 13
> any(is.na(data))
[1] TRUE
> sum(is.na(data))
[1] 149
> ### Having a look at missing data
> md.pattern(data)

```

```

Loan_ID Education ApplicantIncome CoapplicantIncome Property_Area Loan_St
atus
480      1      1      1      1      1
1
43      1      1      1      1      1
1
25      1      1      1      1      1
1
5       1      1      1      1      1
1
19      1      1      1      1      1
1
1       1      1      1      1      1
1
10      1      1      1      1      1
1
1       1      1      1      1      1
1
12      1      1      1      1      1
1
1       1      1      1      1      1
1
1       1      1      1      1      1
1
12      1      1      1      1      1
1
1       1      1      1      1      1
1
1       1      1      1      1      1
1
1       1      1      1      1      1
1
1       0      0      0      0      0
0
Married Gender Loan_Amount_Term Dependents LoanAmount Self_Employed

```

480	1	1	1	1	1	1
43	1	1	1	1	1	1
25	1	1	1	1	1	0
5	1	1	1	1	1	0
19	1	1	1	1	0	1
1	1	1	1	1	0	0
10	1	1	1	1	0	1
1	1	1	1	0	0	1
12	1	1	0	1	1	1
1	1	1	0	1	1	0
1	1	1	0	0	1	1
12	1	0	1	1	1	1
1	1	0	1	1	1	1
2	0	1	1	0	1	1
1	0	1	1	0	0	1
	3	13	14	15	22	32

	Credit_History	
480	1	0
43	0	1
25	1	1
5	0	2
19	1	1
1	0	3
10	1	1
1	1	2
12	1	1
1	1	2
1	1	2
12	1	1
1	0	2
2	1	2
1	1	3
	50	149

```
> aggr_plot <- aggr(data, col=c('Blue','Yellow'), numbers=TRUE, sortVars=TRUE
, labels=names(data), cex.axis=.5,cex.numbers=.9, gap=1, ylab=c("Histogram of
missing data","Pattern"))
```

Variables sorted by number of missings:

Variable	Count
Credit_History	0.081433225
Self_Employed	0.052117264
LoanAmount	0.035830619
Dependents	0.024429967
Loan_Amount_Term	0.022801303
Gender	0.021172638
Married	0.004885993
Loan_ID	0.000000000
Education	0.000000000
ApplicantIncome	0.000000000
CoapplicantIncome	0.000000000
Property_Area	0.000000000
Loan_Status	0.000000000

```
> # dropping the ID column
> Loan_ID = data$Loan_ID
> data$Loan_ID = NULL
> data$Credit_History = factor(data$Credit_History)
> # Exploratory Analysis (Univariate)
> p1= qplot(Gender,data = data,geom="auto")
> p2 = qplot(Married, data=data,geom="auto")
> p3 = qplot(Dependents,data = data,geom="auto")
> p4 = qplot(Education,data = data,geom="auto")
> p5 = qplot(Self_Employed,data = data,geom="auto")
> p6 = qplot(Credit_History,data = data,geom="auto")
> p7 = qplot(Property_Area,data = data,geom="auto")
```

```
> p8 = qplot(Loan_Status,data = data,geom="auto")
> grid.arrange(p1,p2,p3,p4,p5,p6,p7,p8,nrow=3,ncol=3)
> # Imputation
> NAsubset = data[c("Gender","Married","Dependents","Self_Employed","LoanAmount",
"Loan_Amount_Term","Credit_History")]
> summary(NAsubset)
```

	Gender	Married	Dependents	Self_Employed	LoanAmount	Loan_Amount_Term
Female	:112	No :213	0 :345	No :500	Min. : 9.0	Min. : 12
Male	:489	Yes :398	1 :102	Yes : 82	1st Qu.:100.0	1st Qu.:360
NA's	: 13	NA's: 3	2 :101	NA's: 32	Median :128.0	Median :360
			3+ : 51		Mean :146.4	Mean :342
			NA's: 15		3rd Qu.:168.0	3rd Qu.:360
					Max. :700.0	Max. :480
					NA's :22	NA's :14

```
Credit_History
0 : 89
1 :475
NA's: 50
```

```
> set.seed(108)
> imputed=complete(mice(NAsubset))
```

iter	imp	variable				
1	1	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
1	2	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
1	3	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
1	4	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
1	5	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
2	1	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
2	2	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
2	3	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
2	4	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
2	5	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
3	1	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
3	2	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
3	3	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
3	4	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
3	5	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
4	1	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
4	2	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term
4	3	Gender	Married	Dependents	Self_Employed	LoanAmount
Term		Credit_History				Loan_Amount_Term

```

4 4 Gender Married Dependents Self_Employed LoanAmount Loan_Amount_
Term Credit_History
4 5 Gender Married Dependents Self_Employed LoanAmount Loan_Amount_
Term Credit_History
5 1 Gender Married Dependents Self_Employed LoanAmount Loan_Amount_
Term Credit_History
5 2 Gender Married Dependents Self_Employed LoanAmount Loan_Amount_
Term Credit_History
5 3 Gender Married Dependents Self_Employed LoanAmount Loan_Amount_
Term Credit_History
5 4 Gender Married Dependents Self_Employed LoanAmount Loan_Amount_
Term Credit_History
5 5 Gender Married Dependents Self_Employed LoanAmount Loan_Amount_
Term Credit_History

```

```

> data$Gender=imputed$Gender
> data$Married=imputed$Married
> data$Dependents=imputed$Dependents
> data$Self_Employed=imputed$Self_Employed
> data$LoanAmount=imputed$LoanAmount
> data$Loan_Amount_Term=imputed$Loan_Amount_Term
> data$Credit_History=imputed$Credit_History
> summary(data)

```

```

      Gender      Married      Dependents      Education      Self_Employed Applicant
Income
Female:116      No :213      0 :355      Graduate      :480      No :527      Min.   :
150
Male   :498      Yes:401      1 :103      Not Graduate:134      Yes: 87      1st Qu.:
2878
                        2 :104
3812
                        3+: 52
5403
5795
                        Max.   :8
1000
CoapplicantIncome      LoanAmount      Loan_Amount_Term      Credit_History      Property
_Area
Min.   :      0      Min.   : 9.0      Min.   : 12      0: 93      Rural   :
179
1st Qu.:      0      1st Qu.:100.0      1st Qu.:360      1:521      Semiurban:
233
Median : 1188      Median :128.0      Median :360
202
Mean   : 1621      Mean   :146.9      Mean   :342
3rd Qu.: 2297      3rd Qu.:168.0      3rd Qu.:360
Max.   :41667      Max.   :700.0      Max.   :480
Loan_Status
N:192
Y:422

```

```

> plottable1=table(data$Loan_Status,data$Property_Area)
> barplot(plottable1, main="Loan Approval vs Property Area", xlab="Property A
rea",col=c("LightBlue","LightYellow"),legend=rownames(plottable1),beside = TR
UE)
> plottable2=table(data$Loan_Status,data$Credit_History)
> barplot(plottable2, main="Loan Approval vs Credit History", xlab="Credit Hi
story",col=c("LightBlue","LightYellow"),legend=rownames(plottable2),beside =
TRUE)
> plottable3=table(data$Loan_Status,data$Loan_Amount_Term)

```

```

> barplot(plottable3, main="Loan Approval vs Loan Amount term", xlab="Loan amount term", col=c("LightBlue", "LightYellow"), legend=rownames(plottable3), beside = TRUE)
> plottable4=table(data$Loan_Status, data$Self_Employed)
> barplot(plottable4, main="Loan Approval vs Self Employed", xlab="Self_Employed", col=c("LightBlue", "LightYellow"), legend=rownames(plottable4), beside = TRUE)
> plottable5=table(data$Loan_Status, data$Education)
> barplot(plottable5, main="Loan Approval vs Education", xlab="Education", col=c("LightBlue", "LightYellow"), legend=rownames(plottable5), beside = TRUE)
> plottable6=table(data$Loan_Status, data$Dependents)
> barplot(plottable6, main="Loan Approval vs Dependents", xlab="Dependents", col=c("LightBlue", "LightYellow"), legend=rownames(plottable6), beside = TRUE)
> plottable7=table(data$Loan_Status, data$Married)
> barplot(plottable7, main="Loan Approval vs Marital Status", xlab="Marriage", col=c("LightBlue", "LightYellow"), legend=rownames(plottable7), beside = TRUE)
> plottable8=table(data$Loan_Status, data$Gender)
> barplot(plottable8, main="Loan Approval vs Gender", xlab="Gender", col=c("LightBlue", "LightYellow"), legend=rownames(plottable8), beside = TRUE)
> # Correlation Analysis
> numeric_features= data[c("ApplicantIncome", "CoapplicantIncome", "LoanAmount", "Loan_Amount_Term")]
> corTable=cor(numeric_features)
> corTable

```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term
ApplicantIncome	1.00000000	-0.11660458	0.56421656	-0.04864677
CoapplicantIncome	-0.11660458	1.00000000	0.17931270	-0.06471335
LoanAmount	0.56421656	0.17931270	1.00000000	0.04063987
Loan_Amount_Term	-0.04864677	-0.06471335	0.04063987	1.00000000

```

> corrplot( cor(as.matrix(numeric_features), method = "pearson", use = "complete.obs"), is.corr = FALSE, type = "lower", order = "hclust", tl.col = "black", tl.srt = 360)
> # applicant income and loan amount correlated
> # Feature Engineering
> # Add a new feature has a coapplicant
> coAppIn=data$CoapplicantIncome
> for(i in data$CoapplicantIncome){
+   data$CoapplicantIncome[data$CoapplicantIncome!=0.00] = 1.00
+ }
> data$Coapplicant= as.factor(data$CoapplicantIncome)
> # data$CoapplicantIncome= coAppIn
> ## Training & Testing Set
> set.seed(108)
> split=sample.split(data$Loan_Status, SplitRatio = .7)
> train=subset(data, split==T)
> test=subset(data, split==F)
> ## Model Building & CV Using GLM
> Status=glm(Loan_Status~Married+LoanAmount+Credit_History+Property_Area, data=train, family="binomial")
> predGlm=predict(Status, type="response", newdata=test)
> summary(Status)

```

```

Call:
glm(formula = Loan_Status ~ Married + LoanAmount + Credit_History + Property_Area, family = "binomial", data = train)

```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.1290	-0.3714	0.5506	0.7211	2.3792

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-3.042930	0.573079	-5.310	1.10e-07	***
MarriedYes	0.630318	0.261820	2.407	0.01606	*
LoanAmount	-0.002024	0.001290	-1.570	0.11652	
Credit_History1	3.853596	0.492311	7.828	4.97e-15	***
Property_AreaSemiurban	0.857523	0.312383	2.745	0.00605	**
Property_AreaUrban	0.429623	0.303043	1.418	0.15628	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 532.79 on 428 degrees of freedom  
Residual deviance: 396.57 on 423 degrees of freedom  
AIC: 408.57

Number of Fisher Scoring iterations: 5

```
> # Computing accuracy  
> table(test$Loan_Status,predGlm>.5)
```

```
      FALSE TRUE  
N       27   31  
Y        4  123  
> (27+123)/(27+123+31+4)  
[1] 0.8108108
```

```
> set.seed(108)  
> # Decision Tree Model  
> numFolds = trainControl( method = "cv", number = 10 )  
> cpGrid = expand.grid( .cp = seq(0.01,0.5,0.01))  
> train(Loan_Status~Married+LoanAmount+Credit_History+Property_Area,data=train,  
method="rpart",trControl=numFolds,tuneGrid=cpGrid)  
CART
```

```
429 samples  
4 predictor  
2 classes: 'N', 'Y'
```

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 386, 386, 385, 386, 386, 387, ...

Resampling results across tuning parameters:

cp	Accuracy	Kappa
0.01	0.8088719	0.4785112
0.02	0.8088719	0.4785112
0.03	0.8088719	0.4785112
0.04	0.8088719	0.4785112
0.05	0.8088719	0.4785112
0.06	0.8088719	0.4785112
0.07	0.8088719	0.4785112
0.08	0.8088719	0.4785112
0.09	0.8088719	0.4785112
0.10	0.8088719	0.4785112
0.11	0.8088719	0.4785112
0.12	0.8088719	0.4785112
0.13	0.8088719	0.4785112
0.14	0.8088719	0.4785112
0.15	0.8088719	0.4785112
0.16	0.8088719	0.4785112
0.17	0.8088719	0.4785112
0.18	0.8088719	0.4785112

0.19	0.8088719	0.4785112
0.20	0.8088719	0.4785112
0.21	0.8088719	0.4785112
0.22	0.8088719	0.4785112
0.23	0.8088719	0.4785112
0.24	0.8088719	0.4785112
0.25	0.8088719	0.4785112
0.26	0.8088719	0.4785112
0.27	0.8088719	0.4785112
0.28	0.8088719	0.4785112
0.29	0.8088719	0.4785112
0.30	0.8088719	0.4785112
0.31	0.8088719	0.4785112
0.32	0.8088719	0.4785112
0.33	0.8088719	0.4785112
0.34	0.8088719	0.4785112
0.35	0.8088719	0.4785112
0.36	0.8088719	0.4785112
0.37	0.8088719	0.4785112
0.38	0.7929629	0.4208189
0.39	0.7293416	0.1653267
0.40	0.6876925	0.0000000
0.41	0.6876925	0.0000000
0.42	0.6876925	0.0000000
0.43	0.6876925	0.0000000
0.44	0.6876925	0.0000000
0.45	0.6876925	0.0000000
0.46	0.6876925	0.0000000
0.47	0.6876925	0.0000000
0.48	0.6876925	0.0000000
0.49	0.6876925	0.0000000
0.50	0.6876925	0.0000000

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was  $cp = 0.37$ .

```
> decisionTreeModel=rpart(Loan_Status~Married+LoanAmount+Credit_History+Property_Area,data=train,method="class",cp=.37)
> rpart.plot(decisionTreeModel,extra=104, box.palette="GnBu",branch.lty=3, shadow.col="gray", nn=TRUE)
> ## CV with rpart
> predDT=predict(decisionTreeModel,newdata = test,type = "class")
> table(test$Loan_Status,predDT)
  predDT
      N      Y
N    27    31
Y     4   123
> # Accuracy
> (27+123)/(27+123+31+4)
[1] 0.8108108
> # RF Model
> set.seed(108)
> randomForestModel=randomForest(Loan_Status~Married+LoanAmount+Credit_History+Property_Area,data=train,ntree=50,nodesize=10)
> predictRF=predict(randomForestModel,newdata=test)
> table(test$Loan_Status,predictRF)
  predictRF
      N      Y
N    28    30
Y     4   123
> # Accuracy
> (29+122)/(29+122+29+5)
[1] 0.8162162
> # AUC Calculation
> glm_ROC=predict(Status,test,type="response")
```

```

> pred_glm=prediction(glm_ROC,test$Loan_Status)
> perf_glm=performance(pred_glm,"tpr","fpr")
> dt_ROC=predict(decisionTreeModel,test)
> pred_dt=prediction(dt_ROC[,2],test$Loan_Status)
> perf_dt=performance(pred_dt,"tpr","fpr")
> RF_ROC=predict(randomForestModel,test,type="prob")
> pred_RF=prediction(RF_ROC[,2],test$Loan_Status)
> perf_RF=performance(pred_RF,"tpr","fpr")
> auc_glm <- performance(pred_glm,"auc")
> auc_glm <- round(as.numeric(auc_glm@y.values),3)
> auc_dt <- performance(pred_dt,"auc")
> auc_dt <- round(as.numeric(auc_dt@y.values),3)
> auc_RF <- performance(pred_RF,"auc")
> auc_RF <- round(as.numeric(auc_RF@y.values),3)
> print(paste('AUC of Logistic Regression:',auc_glm))
[1] "AUC of Logistic Regression: 0.787"
> print(paste('AUC of Decision Tree:',auc_dt))
[1] "AUC of Decision Tree: 0.717"
> print(paste('AUC of Random Forest:',auc_RF))
[1] "AUC of Random Forest: 0.757"
> # ROC Curves
> plot(perf_glm, main = "ROC curves for the models", col='blue')
> plot(perf_dt,add=TRUE, col='red')
> plot(perf_RF, add=TRUE, col='green3')
> legend('bottom', c("Logistic Regression", "Decision Tree", "Random Forest")
, fill = c('blue','red','green3'), bty='n')
>

```